

# Importing Important Libraries

## Steps To Be Followed

1. Importing necessary Libraries
2. Creating S3 bucket
3. Mapping train And Test Data in S3
4. Mapping The path of the models in S3

```
In [10]: import sagemaker
import boto3
from sagemaker.amazon.amazon_estimator import get_image_uri
from sagemaker.session import s3_input, Session
```

```
In [11]: bucket_name = 'bankapplication' # <--- CHANGE THIS VARIABLE TO A UNIQUE NAME FOR YOUR
my_region = boto3.session.Session().region_name # set the region of the instance
print(my_region)
```

us-east-1

```
In [12]: s3 = boto3.resource('s3')
try:
    if my_region == 'us-east-1':
        s3.create_bucket(Bucket=bucket_name)
        print('S3 bucket created successfully')
except Exception as e:
    print('S3 error: ',e)
```

S3 bucket created successfully

```
In [13]: # set an output path where the trained model will be saved
prefix = 'xgboost-as-a-built-in-algo'
output_path = 's3://{}/{}/output'.format(bucket_name, prefix)
print(output_path)
```

s3://bankapplication/xgboost-as-a-built-in-algo/output

## Downloading The Dataset And Storing in S3

```
In [14]: import pandas as pd
import urllib
try:
    urllib.request.urlretrieve ("https://d1.awsstatic.com/tmt/build-train-deploy-mach
    print('Success: downloaded bank_clean.csv.')
except Exception as e:
    print('Data load error: ',e)

try:
    model_data = pd.read_csv('./bank_clean.csv',index_col=0)
```

```

    print('Success: Data loaded into dataframe.')
except Exception as e:
    print('Data load error: ',e)

```

Success: downloaded bank\_clean.csv.  
 Success: Data loaded into dataframe.

In [7]:

```

### Train Test split

import numpy as np
train_data, test_data = np.split(model_data.sample(frac=1, random_state=1729), [int(0
print(train_data.shape, test_data.shape)

```

(28831, 61) (12357, 61)

In [15]:

```

### Saving Train And Test Into Buckets
## We start with Train Data
import os
pd.concat([train_data['y_yes'], train_data.drop(['y_no', 'y_yes'],
                                                axis=1)],
          axis=1).to_csv('train.csv', index=False)
boto3.Session().resource('s3').Bucket(bucket_name).Object(os.path.join(prefix, 'train/
s3_input_train = sagemaker.s3_input(s3_data='s3://{}/{}/train'.format(bucket_name, pr

```

's3\_input' class will be renamed to 'TrainingInput' in SageMaker Python SDK v2.

In [16]:

```

# Test Data Into Buckets
pd.concat([test_data['y_yes'], test_data.drop(['y_no', 'y_yes'], axis=1)], axis=1).to
boto3.Session().resource('s3').Bucket(bucket_name).Object(os.path.join(prefix, 'test/
s3_input_test = sagemaker.s3_input(s3_data='s3://{}/{}/test'.format(bucket_name, pref

```

's3\_input' class will be renamed to 'TrainingInput' in SageMaker Python SDK v2.

## Building Models Xgboost- Inbuilt Algorithm

In [18]:

```

# this line automatically looks for the XGBoost image URI and builds an XGBoost conta
# specify the repo_version depending on your preference.
container = get_image_uri(boto3.Session().region_name,
                          'xgboost',
                          repo_version='1.0-1')

```

'get\_image\_uri' method will be deprecated in favor of 'ImageURIProvider' class in SageMaker Python SDK v2.

In [25]:

```

# initialize hyperparameters
hyperparameters = {
    "max_depth": "5",
    "eta": "0.2",

```

```
"gamma": "4",  
"min_child_weight": "6",  
"subsample": "0.7",  
"objective": "binary:logistic",  
"num_round": 50  
}
```

```
In [26]: # construct a SageMaker estimator that calls the xgboost-container  
estimator = sagemaker.estimator.Estimator(image_name=container,  
                                           hyperparameters=hyperparameters,  
                                           role=sagemaker.get_execution_role(),  
                                           train_instance_count=1,  
                                           train_instance_type='ml.m5.2xlarge',  
                                           train_volume_size=5, # 5 GB  
                                           output_path=output_path,  
                                           train_use_spot_instances=True,  
                                           train_max_run=300,  
                                           train_max_wait=600)
```

Parameter image\_name will be renamed to image\_uri in SageMaker Python SDK v2.

```
In [27]: estimator.fit({'train': s3_input_train, 'validation': s3_input_test})
```

```

2020-08-29 09:49:29 Starting - Starting the training job...
2020-08-29 09:49:31 Starting - Launching requested ML instances.....
2020-08-29 09:51:04 Starting - Preparing the instances for training...
2020-08-29 09:51:44 Downloading - Downloading input data
2020-08-29 09:51:44 Training - Downloading the training image..INFO:sagemaker-container
s:Imported framework sagemaker_xgboost_container.training
INFO:sagemaker-containers:Failed to parse hyperparameter objective value binary:logisti
c to Json.
Returning the value itself
INFO:sagemaker-containers:No GPUs detected (normal if no gpus installed)
INFO:sagemaker_xgboost_container.training:Running XGBoost Sagemaker in algorithm mode
INFO:root:Determined delimiter of CSV input is ','
INFO:root:Determined delimiter of CSV input is ','
INFO:root:Determined delimiter of CSV input is ','
[09:52:07] 28831x59 matrix with 1701029 entries loaded from /opt/ml/input/data/train?fo
rmat=csv&label_column=0&delimiter=,
INFO:root:Determined delimiter of CSV input is ','
[09:52:07] 12357x59 matrix with 729063 entries loaded from /opt/ml/input/data/validatio
n?format=csv&label_column=0&delimiter=,
INFO:root:Single node training.
INFO:root:Train matrix has 28831 rows
INFO:root:Validation matrix has 12357 rows
[09:52:07] WARNING: /workspace/src/learner.cc:328:
Parameters: { num_round } might not be used.

```

This may not be accurate due to some parameters are only used in language bindings bu  
t  
passed down to XGBoost core. Or some parameters are not used but slip through this  
verification. Please open an issue if you find above cases.

```

[0]#011train-error:0.10079#011validation-error:0.10528
[1]#011train-error:0.09968#011validation-error:0.10456
[2]#011train-error:0.10017#011validation-error:0.10375
[3]#011train-error:0.09989#011validation-error:0.10310
[4]#011train-error:0.09996#011validation-error:0.10286
[5]#011train-error:0.09906#011validation-error:0.10261
[6]#011train-error:0.09930#011validation-error:0.10286
[7]#011train-error:0.09951#011validation-error:0.10261
[8]#011train-error:0.09920#011validation-error:0.10286
[9]#011train-error:0.09871#011validation-error:0.10294
[10]#011train-error:0.09868#011validation-error:0.10294
[11]#011train-error:0.09868#011validation-error:0.10326
[12]#011train-error:0.09854#011validation-error:0.10358
[13]#011train-error:0.09892#011validation-error:0.10342
[14]#011train-error:0.09850#011validation-error:0.10342
[15]#011train-error:0.09844#011validation-error:0.10326
[16]#011train-error:0.09857#011validation-error:0.10318
[17]#011train-error:0.09799#011validation-error:0.10318
[18]#011train-error:0.09816#011validation-error:0.10383
[19]#011train-error:0.09857#011validation-error:0.10383
[20]#011train-error:0.09830#011validation-error:0.10350
[21]#011train-error:0.09826#011validation-error:0.10318
[22]#011train-error:0.09847#011validation-error:0.10399
[23]#011train-error:0.09833#011validation-error:0.10407
[24]#011train-error:0.09812#011validation-error:0.10415
[25]#011train-error:0.09812#011validation-error:0.10399
[26]#011train-error:0.09774#011validation-error:0.10375
[27]#011train-error:0.09781#011validation-error:0.10375
[28]#011train-error:0.09781#011validation-error:0.10391

```

```
[29]#011train-error:0.09778#011validation-error:0.10367
[30]#011train-error:0.09781#011validation-error:0.10383
[31]#011train-error:0.09771#011validation-error:0.10358
[32]#011train-error:0.09743#011validation-error:0.10391
[33]#011train-error:0.09753#011validation-error:0.10342
[34]#011train-error:0.09767#011validation-error:0.10342
[35]#011train-error:0.09757#011validation-error:0.10350
[36]#011train-error:0.09757#011validation-error:0.10342
[37]#011train-error:0.09736#011validation-error:0.10342
[38]#011train-error:0.09750#011validation-error:0.10342
[39]#011train-error:0.09733#011validation-error:0.10350
[40]#011train-error:0.09705#011validation-error:0.10358
[41]#011train-error:0.09701#011validation-error:0.10383
[42]#011train-error:0.09712#011validation-error:0.10407
[43]#011train-error:0.09698#011validation-error:0.10375
[44]#011train-error:0.09733#011validation-error:0.10342
[45]#011train-error:0.09736#011validation-error:0.10367
[46]#011train-error:0.09746#011validation-error:0.10350
[47]#011train-error:0.09736#011validation-error:0.10358
[48]#011train-error:0.09712#011validation-error:0.10334
[49]#011train-error:0.09712#011validation-error:0.10318
```

2020-08-29 09:52:19 Uploading - Uploading generated training model

2020-08-29 09:52:19 Completed - Training job completed

Training seconds: 44

Billable seconds: 21

Managed Spot Training savings: 52.3%

## Deploy Machine Learning Model As Endpoints

In [28]: `xgb_predictor = estimator.deploy(initial_instance_count=1,instance_type='ml.m4.xlarge`

Parameter image will be renamed to image\_uri in SageMaker Python SDK v2.  
-----!

## Prediction of the Test Data

In [29]: `from sagemaker.predictor import csv_serializer
test_data_array = test_data.drop(['y_no', 'y_yes'], axis=1).values #load the data into
xgb_predictor.content_type = 'text/csv' # set the data type for an inference
xgb_predictor.serializer = csv_serializer # set the serializer type
predictions = xgb_predictor.predict(test_data_array).decode('utf-8') # predict!
predictions_array = np.fromstring(predictions[1:], sep=',') # and turn the prediction
print(predictions_array.shape)`

(12357,)

In [30]: `predictions_array`

Out[30]: `array([0.05214286, 0.05660191, 0.05096195, ..., 0.03436061, 0.02942475,
0.03715819])`

In [31]: `cm = pd.crosstab(index=test_data['y_yes'], columns=np.round(predictions_array), rowna`

```

tn = cm.iloc[0,0]; fn = cm.iloc[1,0]; tp = cm.iloc[1,1]; fp = cm.iloc[0,1]; p = (tp+tn)/cm.sum(axis=0)[1]
print("\n{0:<20}{1:<4.1f}%\n".format("Overall Classification Rate: ", p))
print("{0:<15}{1:<15}{2:>8}".format("Predicted", "No Purchase", "Purchase"))
print("Observed")
print("{0:<15}{1:<2.0f}% ({2:<}){3:>6.0f}% ({4:<})".format("No Purchase", tn/(tn+fn)*100, "No Purchase", tn, "No Purchase"))
print("{0:<16}{1:<1.0f}% ({2:<}){3:>7.0f}% ({4:<}) \n".format("Purchase", fn/(tn+fn)*100, "Purchase", fn, "Purchase"))

```

Overall Classification Rate: 89.7%

Predicted	No Purchase	Purchase
Observed		
No Purchase	91% (10785)	34% (151)
Purchase	9% (1124)	66% (297)

## Deleting The Endpoints

In [32]:

```

sagemaker.Session().delete_endpoint(xgb_predictor.endpoint)
bucket_to_delete = boto3.resource('s3').Bucket(bucket_name)
bucket_to_delete.objects.all().delete()

```

Out[32]:

```

[{'ResponseMetadata': {'RequestId': '2FF829102DC6DFD1',
  'HostId': 'mYPqeWyx3REoLIIsQu2MVorzKLrlxES2n6Dcdr3PycVf1VkrCxicEewoPP8IxrGuc5MGksLnjynY=',
  'HTTPStatusCode': 200,
  'HTTPHeaders': {'x-amz-id-2': 'mYPqeWyx3REoLIIsQu2MVorzKLrlxES2n6Dcdr3PycVf1VkrCxicEewoPP8IxrGuc5MGksLnjynY=',
    'x-amz-request-id': '2FF829102DC6DFD1',
    'date': 'Sat, 29 Aug 2020 10:21:27 GMT',
    'connection': 'close',
    'content-type': 'application/xml',
    'transfer-encoding': 'chunked',
    'server': 'AmazonS3'},
  'RetryAttempts': 0},
  'Deleted': [{'Key': 'xgboost-as-a-built-in-algo/train/train.csv'},
    {'Key': 'xgboost-as-a-built-in-algo/test/test.csv'},
    {'Key': 'xgboost-as-a-built-in-algo/output/sagemaker-xgboost-2020-08-29-09-49-29-015/output/model.tar.gz'}]]

```

In [ ]: